

Experiment 8: Implement the Reinforcement Learning with OpenAI Gym using the FrozenLake Environment

Name: Nabeel Shan - BESE-29C

```
# !pip install gymnasium matplotlib numpy
```

```
# Imports and Environment Registration
import gymnasium as gym
from gymnasium.envs.registration import register
import numpy as np
import matplotlib.pyplot as plt
import random

# 1. Safely register custom non-slippery environments
try:
    register(
        id='FrozenLakeNotSlippery-v1',
        entry_point='gymnasium.envs.toy_text.frozen_lake:FrozenLakeEnv',
        kwargs={'map_name' : '4x4', 'is_slippery': False},
        max_episode_steps=100,
        reward_threshold=0.8196,
    )
    register(
        id='FrozenLake8x8NotSlippery-v1',
        entry_point='gymnasium.envs.toy_text.frozen_lake:FrozenLakeEnv',
        kwargs={'map_name' : '8x8', 'is_slippery': False},
        max_episode_steps=100,
        reward_threshold=0.8196,
    )
    print("Environments registered successfully!")
except Exception as e:
    print(f>Note: {e}")

# Initialize the 4x4 non-slippery environment
env = gym.make('FrozenLakeNotSlippery-v1')

print(f"Observation Space: {env.observation_space.n} states")
print(f"Action Space: {env.action_space.n} actions")

# Initialize the Q-Table with zeros
Q = np.zeros([env.observation_space.n, env.action_space.n])
```

```
Environments registered successfully!
Observation Space: 16 states
Action Space: 4 actions
```

```
# Core Reinforcement Learning Functions

# Hyperparameters
alpha = 0.8          # Learning rate: How much new info overrides old info
gamma = 0.95         # Discount factor: Importance of future rewards
epsilon = 1.0        # Exploration rate: Start at 100% random actions
epsilon_decay = 0.995
min_epsilon = 0.01

def choose_action(state, current_epsilon):
    """
    Epsilon-Greedy Action Selection.
    """
    # Explore: take a random action
    if random.uniform(0, 1) < current_epsilon:
        return env.action_space.sample()
    # Exploit: take the action with the highest Q-value for this state
    else:
        return np.argmax(Q[state, :])

def learn(s, s1, r, a):
    """
    Q-Learning Update Rule (Bellman Equation).
    """
    # Maximum future reward for the next state
    max_future_q = np.max(Q[s1, :])
```

```
# Update the Q-value for the current state and action
Q[s, a] = Q[s, a] + alpha * (r + gamma * max_future_q - Q[s, a])
```

```
# Training the Agent

num_episodes = 2000
max_steps_per_episode = 100
rewards_all_episodes = []

print("Training Agent...")

for i in range(num_episodes):
    # Modern reset returns a tuple: (state, info)
    s, info = env.reset()

    d = False
    rewards_current_episode = 0

    for step in range(max_steps_per_episode):
        # 1. Choose action
        a = choose_action(s, epsilon)

        # 2. Take action (Modern syntax returns 5 values)
        s1, r, terminated, truncated, _ = env.step(a)

        # 3. Update Q-Table
        learn(s, s1, r, a)

        # 4. Check if episode is over
        d = terminated or truncated
        rewards_current_episode += r

        if d:
            break

        # 5. Move to next state
        s = s1

    # Decay epsilon after each episode so the agent explores less over time
    epsilon = max(min_epsilon, epsilon * epsilon_decay)

    rewards_all_episodes.append(rewards_current_episode)

print("Training Complete!")
print(f"Final Epsilon (Exploration Rate): {epsilon:.4f}")
```

```
Training Agent...
Training Complete!
Final Epsilon (Exploration Rate): 0.0100
```

```
# Analytics and Visualization


# Calculate moving average of rewards
window = 100
moving_avg = np.convolve(rewards_all_episodes, np.ones(window)/window, mode='valid')

# Plot the learning curve
plt.figure(figsize=(10, 5))
plt.plot(moving_avg, color='blue', linewidth=2)
plt.title('Agent Learning Curve (100-Episode Moving Average)')
plt.xlabel('Episode')
plt.ylabel('Success Rate (Goal Reached)')
plt.grid(True)
plt.show()

print("\n--- Final Q-Table (The Agent's Brain) ---")
print("Rows = States (0-15) | Columns = Actions (Left, Down, Right, Up)")
print(np.round(Q, 3))

# Print the optimal policy mapped to human-readable directions
actions_dict = {0: 'Left', 1: 'Down', 2: 'Right', 3: 'Up'}
print("\n--- Optimal Policy Extracted ---")
policy = []
for state in range(env.observation_space.n):
    # If a state is all zeros, it's a hole or the goal (terminal state)
    if np.max(Q[state]) > 0:
        best_action = np.argmax(Q[state])
        policy.append(actions_dict[best_action])
    else:
        policy.append("TERM")
```

```
# Reshape policy into 4x4 grid
policy_grid = np.array(policy).reshape(4, 4)
print(policy_grid)
```



```
--- Final Q-Table (The Agent's Brain) ---
Rows = States (0-15) | Columns = Actions (Left, Down, Right, Up)
[[0.735 0.774 0.698 0.735]
 [0.735 0.    0.66  0.559]
 [0.698 0.    0.    0.53 ]
 [0.    0.    0.    0.    ]
 [0.774 0.815 0.    0.735]
 [0.    0.    0.    0.    ]
 [0.    0.    0.    0.    ]
 [0.    0.    0.    0.    ]
 [0.815 0.    0.857 0.774]
 [0.815 0.902 0.902 0.    ]
 [0.841 0.95  0.    0.    ]
 [0.    0.    0.    0.    ]
 [0.    0.    0.    0.    ]
 [0.    0.902 0.95  0.857]
 [0.902 0.95  1.    0.902]
 [0.    0.    0.    0.    ]]

--- Optimal Policy Extracted ---
[['Down' 'Left' 'Left' 'TERM']
 ['Down' 'TERM' 'TERM' 'TERM']
 ['Right' 'Down' 'Down' 'TERM']
 ['TERM' 'Right' 'Right' 'TERM']]
```

Start coding or [generate](#) with AI.

Task 1 - Wins & Losses Performance Summary

```
# Track Wins & Losses Summary

total_episodes = len(rewards_all_episodes)
total_wins = int(sum(rewards_all_episodes))
total_losses = total_episodes - total_wins
overall_win_rate = (total_wins / total_episodes) * 100

print("=====")
print("    RL AGENT PERFORMANCE SUMMARY    ")
print("=====")
print(f"Total Episodes Played : {total_episodes}")
print(f"Total Wins (Goal)      : {total_wins}")
print(f"Total Losses (Hole)    : {total_losses}")
print(f"Overall Win Rate       : {overall_win_rate:.2f}%")
print("=====")

# A+ Engineering Addition: Segmented Evaluation
# Proving that the agent got better over time by comparing the
# first 25% of training (high exploration) to the last 25% (high exploitation).

quarter_mark = total_episodes // 4
first_quarter_wins = sum(rewards_all_episodes[:quarter_mark])
```

```

last_quarter_wins = sum(rewards_all_episodes[-quarter_mark:])

first_quarter_rate = (first_quarter_wins / quarter_mark) * 100
last_quarter_rate = (last_quarter_wins / quarter_mark) * 100

print("\nLEARNING PROGRESSION (Convergence Proof):")
print(f"Phase 1 (First 25% - High Randomness) : {first_quarter_rate:.2f}% Win Rate")
print(f"Phase 4 (Last 25% - Exploiting Q-Table): {last_quarter_rate:.2f}% Win Rate")

if last_quarter_rate > first_quarter_rate:
    print("\nConclusion: The agent successfully converged on an optimal policy.")
else:
    print("\nConclusion: The agent failed to learn. Adjust hyperparameters (alpha, gamma).")

```

```

=====
      RL AGENT PERFORMANCE SUMMARY
=====
Total Episodes Played : 2000
Total Wins (Goal)      : 1764
Total Losses (Hole)    : 236
Overall Win Rate       : 88.20%
=====

LEARNING PROGRESSION (Convergence Proof):
Phase 1 (First 25% - High Randomness) : 56.60% Win Rate
Phase 4 (Last 25% - Exploiting Q-Table): 99.00% Win Rate

Conclusion: The agent successfully converged on an optimal policy.

```

Start coding or [generate](#) with AI.

Task 2 - The Epsilon-Greedy Strategy

```

import numpy as np
import random

# For demonstration in this specific cell, let's assume a sample state and Q-table
# In your actual loop, these are dynamically updated.
sample_state = 0
sample_epsilon = 0.5 # 50% chance to explore, 50% chance to exploit

def choose_action_epsilon_greedy(state, current_epsilon):
    """
    Task 2 Implementation: Epsilon-Greedy Action Selection
    Replaces the purely random np.random.choice([0,1,2,3])
    """

    # 1. EXPLORATION: Roll a random number between 0 and 1.
    # If it's less than epsilon, the agent acts randomly to discover new paths.
    if random.uniform(0, 1) < current_epsilon:
        action = env.action_space.sample()
        action_type = "Exploration (Random)"

    # 2. EXPLOITATION: Otherwise, use the agent's "brain" (the Q-table).
    # np.argmax finds the action with the highest expected reward for this specific state.
    else:
        action = np.argmax(Q[state, :])
        action_type = "Exploitation (Q-Table)"

    return action, action_type

# --- Demonstration of the strategy ---
print("Testing Epsilon-Greedy Logic:")
for test in range(5):
    selected_action, logic_used = choose_action_epsilon_greedy(sample_state, sample_epsilon)
    print(f"Test {test+1}: Agent chose action {selected_action} using {logic_used}")

```

```

Testing Epsilon-Greedy Logic:
Test 1: Agent chose action 1 using Exploitation (Q-Table)
Test 2: Agent chose action 1 using Exploitation (Q-Table)
Test 3: Agent chose action 1 using Exploitation (Q-Table)
Test 4: Agent chose action 1 using Exploitation (Q-Table)
Test 5: Agent chose action 3 using Exploration (Random)

```

Start coding or [generate](#) with AI.

Task 3 - Q-Learning Implementation

```

# Define the hyperparameters required for the formula
alpha = 0.8 # Learning Rate: How quickly the agent abandons old beliefs for new ones
gamma = 0.95 # Discount Factor: How much the agent cares about long-term vs short-term rewards

def learn_q_table(s, s1, r, a):
    """
    Task 3 Implementation: The Bellman Update Equation.
    This function updates the agent's brain (Q-Table) after every single step.

    Parameters:
    s : Current State
    s1 : Next State
    r : Immediate Reward received
    a : Action taken
    """

    # Step 1: Find the maximum possible future reward we can get from the NEXT state (s1)
    # np.max looks at all 4 possible actions in state s1 and picks the highest Q-value
    max_future_q = np.max(Q[s1, :])

    # Step 2: Calculate the Temporal Difference (TD) Target
    # This is what the agent *actually* experienced plus what it *expects* to get next
    td_target = r + (gamma * max_future_q)

    # Step 3: Calculate the Temporal Difference (TD) Error
    # The difference between what happened (td_target) and what the agent originally thought (Q[s, a])
    td_error = td_target - Q[s, a]

    # Step 4: Update the Q-Value using the learning rate (alpha)
    Q[s, a] = Q[s, a] + (alpha * td_error)

    # Note: Steps 2, 3, and 4 are usually condensed into one line in production code:
    # Q[s, a] = Q[s, a] + alpha * (r + gamma * np.max(Q[s1, :]) - Q[s, a])

    # --- Demonstration to verify the function works ---
    print("Simulating a learning update...")
    print(f"Old Q-Value for State 0, Action 1: {Q[0, 1]}")

    # Simulate: Agent moves from State 0 to State 4 (Action 1), receives 0 reward.
    learn_q_table(s=0, s1=4, r=0, a=1)

    print(f"New Q-Value for State 0, Action 1: {Q[0, 1]}")
    print("Task 3 Complete: The agent's memory has been successfully updated.")

```

```

Simulating a learning update...
Old Q-Value for State 0, Action 1: 0.7737809374999999
New Q-Value for State 0, Action 1: 0.7737809374999999
Task 3 Complete: The agent's memory has been successfully updated.

```

Start coding or [generate](#) with AI.

Task 4 - 4x4 vs 8x8 Performance Comparison

```

import gymnasium as gym
import numpy as np
import matplotlib.pyplot as plt
import random

def train_agent_experiment(env_name, num_episodes=5000, max_steps=200, epsilon_decay=0.999):
    """
    A robust, reusable function to train a Q-learning agent on any given environment.
    Returns the moving average of rewards for plotting.
    """
    env = gym.make(env_name)
    Q = np.zeros([env.observation_space.n, env.action_space.n])

    # Hyperparameters
    alpha = 0.8
    gamma = 0.95
    epsilon = 1.0
    min_epsilon = 0.01

    rewards_all = []

    for i in range(num_episodes):
        s, _ = env.reset()
        d = False
        episode_reward = 0

```

```

    for step in range(max_steps):
        # Task 2 Logic: Epsilon-Greedy
        if random.uniform(0, 1) < epsilon:
            a = env.action_space.sample()
        else:
            a = np.argmax(Q[s, :])

        s1, r, terminated, truncated, _ = env.step(a)
        d = terminated or truncated

        # Task 3 Logic: Q-Learning Update
        max_future_q = np.max(Q[s1, :])
        Q[s, a] = Q[s, a] + alpha * (r + gamma * max_future_q - Q[s, a])

        s = s1
        episode_reward += r
        if d:
            break

    epsilon = max(min_epsilon, epsilon * epsilon_decay)
    rewards_all.append(episode_reward)

# Calculate a 100-episode moving average for smooth plotting
window = 100
moving_avg = np.convolve(rewards_all, np.ones(window)/window, mode='valid')
return moving_avg

print("Starting Lab Experiment: Scale Comparison...")

# 1. Train on the 4x4 Grid
print("Training Agent on 4x4 Grid (State Space: 16)...")
# Note: We use a faster decay for 4x4 because it's easier to solve
rewards_4x4 = train_agent_experiment(
    'FrozenLakeNotSlippery-v1',
    num_episodes=5000,
    max_steps=100,
    epsilon_decay=0.995
)

# 2. Train on the 8x8 Grid
print("Training Agent on 8x8 Grid (State Space: 64)...")
# Note: We use a slower decay for 8x8 so it explores the massive map longer
rewards_8x8 = train_agent_experiment(
    'FrozenLake8x8NotSlippery-v1',
    num_episodes=5000,
    max_steps=200,
    epsilon_decay=0.999
)


print("Training Complete! Generating Analytics Plot...")

# 3. Plot the Comparison
plt.figure(figsize=(12, 6))
plt.plot(rewards_4x4, color='blue', label='4x4 Grid (16 States)', linewidth=2)
plt.plot(rewards_8x8, color='orange', label='8x8 Grid (64 States)', linewidth=2)

plt.title('Agent Performance: Small vs Large Environment Convergence', fontsize=14, fontweight='bold')
plt.xlabel('Episodes', fontsize=12)
plt.ylabel('Success Rate (100-Episode Moving Average)', fontsize=12)
plt.legend(loc='lower right', fontsize=12)
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```

Starting Lab Experiment: Scale Comparison...
Training Agent on 4x4 Grid (State Space: 16)...
Training Agent on 8x8 Grid (State Space: 64)...
Training Complete! Generating Analytics Plot...



```
for reward in rewards_4x4:  
    print(reward)
```

```
0.05  
0.060000000000000005  
0.060000000000000005  
0.07  
0.08  
0.07999999999999999  
0.07999999999999999  
0.09  
0.09999999999999999  
0.10999999999999999  
0.10999999999999999  
0.10999999999999999  
0.09999999999999999  
0.1  
0.1  
0.1  
0.11  
0.11  
0.12  
0.13  
0.13  
0.13  
0.14  
0.15000000000000002  
0.15000000000000002  
0.15000000000000002  
0.15000000000000002  
0.15000000000000002  
0.16000000000000003  
0.16000000000000003  
0.16000000000000003  
0.16000000000000003  
0.16  
0.16  
0.16  
0.17  
0.18000000000000002  
0.19000000000000003  
0.19000000000000003  
0.2  
0.2  
0.21000000000000002  
0.22000000000000003  
0.23000000000000004  
0.23000000000000004  
0.24000000000000002  
0.24000000000000002
```

Start coding or generate with AI.